




This image shows a blank sheet of white paper with horizontal blue or grey ruling lines, typical of notebook paper. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.


-
- This image shows a blank sheet of white paper with horizontal blue ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins or other markings on the paper.

[illegible]

- [illegible]

[illegible]

- [illegible]



[illegible]This image shows a single sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page. There are approximately 20 lines visible. The paper has a slightly textured appearance and is set against a dark background.[illegible]

This image shows a blank sheet of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page. There are approximately 20 lines visible. On the left side, there is a vertical margin line, creating a narrow left margin. The paper appears to be from a notebook or a standard ruled document.[illegible][illegible]

Dilated convolution

* Instead of ordinary convolution (1D):

$$(I * k)(t) = \sum_{\tau} I(t - \tau) k(\tau)$$


Use:

$$(I \times_e k)(t) = \sum_{\tau} I(t - l\tau) k(\tau)$$

take steps of size l in image when performing the convolution


Pooling

- Pooling = down sampling spatial dimensions (depth unchanged)
- Max pooling: partition non overlapping regions and choose max in each region
- Using 2 x 2 regions reduces the layer dimensions by 75%
- Alternatives:
 - Average pooling
 - L2 norm pooling
 - ROI pooling (output size is fixed and input size is variable)



Pooling


- Pooling is done by scanning with a filter with stride
- Stride is normally selected without filter overlap (downsampling). For example a 2 x 2 filter with a stride of 2
- Convolution with a stride can also be used to downsample but this will average and so we normally downsample with pooling instead of convolution
- Pooling retains more information (e.g. indicates if a feature is there or not)



Pooling

- Pulling is a layer without parameters and has no learning
- The depth dimension is normally not pooled
- Pulling supports multiple scale analysis (a 3x3 window in a pooled layer covers a larger area in the layer before it)
- Pulling helps in reducing the number of network coefficients
- The amount of pooling is a design choice (hyperparameters)

Example network:



Keras MNIST example

```
from keras import layers
from keras import models
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.summary()
```

```
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
model.summary()
```

Convolution parameters

Conv2D(output_depth, (window_height, window_width))

strides = 1, padding = 'valid' ('valid': no padding, 'same': yes padding)
dilation-rate = 1

Keras MNIST example

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
maxpooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
maxpooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_3 (Conv2D)	(None, 3, 3, 64)	36928
flatten_1 (Flatten)	(None, 576)	0
dense_1 (Dense)	(None, 64)	36928
dense_2 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

(x) Downsample

(x x) Reduction by 2 due to no zero padding

32 3x3x1 filters
(32 x (3x3x1+1))

64 3x3x32 filters
(64 x (3x3x32+1))

64 3x3x64 filters
(64 x (3x3x64+1))

2x3x64
64 units with
576 inputs
(64 x (576+1))

10 units with
64 inputs
(10 x (64+1))

Keras MNIST example

✧ Train on MNIST:

```
from keras.datasets import mnist
from keras.utils import to_categorical

# load data
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

# compile model, fit, and evaluate
model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(train_images, train_labels, epochs=5, batch_size=64)
test_loss, test_acc = model.evaluate(test_images, test_labels)

test_acc                # 0.99080000000000001
```

Convnets improved the 97.8% accuracy of a fully connected network to 99.1%

Keras cats/dogs example

- Cats and dogs classification (Kaggle challenge)
- We use a subset of 2000 cats + 2000 dogs
- Larger image size (150 x 150) compared with MNIST (28 x 28) and hence there's a need for a deeper network

```
from keras import layers
from keras import models
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.summary()
```

Keras cats/dogs example

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
maxpooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
maxpooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
maxpooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
maxpooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

Keras cats/dogs example

* Compile network:

```
from keras import optimizers
model.compile(
    loss='binary_crossentropy',
    optimizer=optimizers.RMSprop(lr=1e-4),
    metrics=['acc'])
```

* Pre-process data:


- load and rescale to [0,1] using a keras data generator

Keras cats/dogs example

* pre-process data using ImageDataGenerator :

```
from keras.preprocessing.image import ImageDataGenerator
train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)
train_generator = train_datagen.flow_from_directory(
    train_dir,                # image folder
    target_size=(150, 150),   # resize images
    batch_size=20,
    class_mode='binary')
validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
```

- Generate 20x150x150x3 tensors for training and validation and corresponding binary labels



Keras cats/dogs example

* Fit model:

use 'fit_generator' instead of 'fit'

- The # of draws of size 'batch_size' to complete a complete epoch is:

$$\text{Steps-per-epoch} = \frac{\text{\# examples}}{\text{batch-size}} = \frac{2000}{20} = 100$$

```
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=30,
    validation_data=validation_generator,
    validation_steps=50)
model.save('cats_and_dogs_small_1.h5')
```

Keras cats/dogs example


* plot results.

```
import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

# Plot accuracy
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

# Plot loss
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

Keras cats/dogs example



- overfit after 5 iterations due to small dataset
67% validation accuracy

Keras cats/dogs example

* To further prevent overfitting, add dropout:

```
# define a new convnet that includes dropout
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```


Keras cats/dogs example

- To address the problem with limited data:
 - Data augmentation:** add examples with perturbations (e.g., rotation, flip, contrast change)
 - Transfer learning:** use a pre-trained convolution base
- Freeze loaded weights of pre-trained blocks after loading them



Data augmentation

```
train_datagen = ImageDataGenerator(                # different from before
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,)
test_datagen = ImageDataGenerator(rescale=1./255)  # as before


train_generator = train_datagen.flow_from_directory( # as before
    * train_dir,
    * target_size=(150, 150),
    * batch_size=32,
    * class_mode='binary')
validation_generator = test_datagen.flow_from_directory( # as before
    * validation_dir,
    * target_size=(150, 150),
    * batch_size=32,
    * class_mode='binary')

history = model.fit_generator(                      # as before
    * train_generator,
    * steps_per_epoch=100,
    * epochs=100,
    * validation_data=validation_generator,
    * validation_steps=50)

model.save('cats_and_dogs_small_2.h5')             # as before
```

no augmentation
during testing
on validation


Data augmentation



- does not overfit. 82% validation accuracy
- Additional improvement is possible with hyperparameter tuning.

Transfer learning

- Use a pre-trained convnet trained on a large data set (e.g., ImageNet object classification)
- ImageNet: 1.4 million labeled images, 1000 classes (animals, objects)
- Use the convolution layers of VGG16 to extract features for the cat/dog problem (representation learning)
- Whether to use higher convolution layers depends on how similar the data sets are



Transfer learning

* Pre-trained models available in Keras:

Models for image classification with weights trained on ImageNet:

- Xception
- VGG16
- VGG19
- ResNet, ResNetV2, ResNeXt
- InceptionV3
- InceptionResNetV2
- MobileNet
- MobileNetV2
- DenseNet
- NASNet

Transfer learning

* Load VGG16

```
from keras.applications import VGG16
conv_base = VGG16(
    weights='imagenet',          # weights checkpoint from which to initialize model
    include_top=False,          # do not include the fully connected layers
                                # (responsible for classifying 1000 classes)
    input_shape=(150, 150, 3))  # optional
conv_base.summary()
```

Transfer learning

```

>>> conv_base.summary()
Layer (type)                                     Output Shape                                     Param #
=====
input_1 (InputLayer)                             (None, 150, 150, 3)                             0

block1_conv1 (Convolution2D)                     (None, 150, 150, 64)                            1192
block1_pool (MaxPooling2D)                       (None, 75, 75, 64)                              0
block2_conv1 (Convolution2D)                     (None, 75, 75, 128)                             7384
block2_pool (MaxPooling2D)                       (None, 37, 37, 128)                              0
block3_conv1 (Convolution2D)                     (None, 37, 37, 256)                             295144
block3_pool (MaxPooling2D)                       (None, 18, 18, 256)                              0
block4_conv1 (Convolution2D)                     (None, 18, 18, 512)                             1180160
block4_pool (MaxPooling2D)                       (None, 9, 9, 512)                                0
block5_conv1 (Convolution2D)                     (None, 9, 9, 512)                                2158908
block5_pool (MaxPooling2D)                       (None, 4, 4, 512)                                0
Total params: 16,714,688
Trainable params: 16,714,688
Non-trainable params: 0
    
```

* 14M parameters

* The input feature map is 150x150x3

* The final feature map is 4x4x512

Transfer learning

• Add pre-trained layers to the network:

- Add conv-base (the loaded model) as a layer
- Freeze weights of pre-trained network
- Train end-to-end

• Larger and slower network

```

from keras import models
from keras import layers
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.summary()
    
```

Transfer learning

```

Layer (type)                                     Output Shape                                     Param #
=====
vgg16 (Model)                                     (None, 4, 4, 512)                             14714688

flatten_1 (Flatten)                              (None, 8192)                                    0

dense_1 (Dense)                                   (None, 256)                                     2097408


dense_2 (Dense)                                   (None, 1)                                       257
=====
Total params: 16,812,353
Trainable params: 16,812,353
Non-trainable params: 0
    
```

- Fine tuning:
 - After training the fully connected layers, unfreeze some top layers in the conv-base and retrain to allow the model to fit the data
- Steps:
 1. Add custom network on top of the trained layers
 2. Freeze the trained layers
 3. Train the custom network
 4. Unfreeze the top layers in the base network
 5. Jointly train the custom network and unfrozen layers

Transfer learning

conv_base.summary()

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 256, 150, 3)	0
=====		
block1_conv1 (Convolution2D)	(None, 128, 128, 64)	1792
block1_conv2 (Convolution2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Convolution2D)	(None, 128, 128, 128)	73856
block2_conv2 (Convolution2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Convolution2D)	(None, 64, 64, 256)	295168
block3_conv2 (Convolution2D)	(None, 64, 64, 256)	580384
block3_conv3 (Convolution2D)	(None, 64, 64, 256)	580384
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Convolution2D)	(None, 16, 16, 512)	1140160
block4_conv2 (Convolution2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Convolution2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
=====		
block5_conv1 (Convolution2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Convolution2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Convolution2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params:	14714880	




Transfer learning

- Freezing all layers up to block number 5:
 - Step through all the layers
 - Set “trainable” starting with block5_conv1

```
conv_base.trainable = True
set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False
```

Transfer learning



* Smoothed curves: 97% accuracy (1% improvement)
 * final test accuracy:

```
test_loss, test_acc = model.evaluate_generator(
    test_generator,
    steps=50)


print('test acc:', test_acc) # 97%
```

[illegible]

CNN architectures

* Google Net / Inception (2014): 22 layers + ^{fewer} parameters (5M)

* Inception block

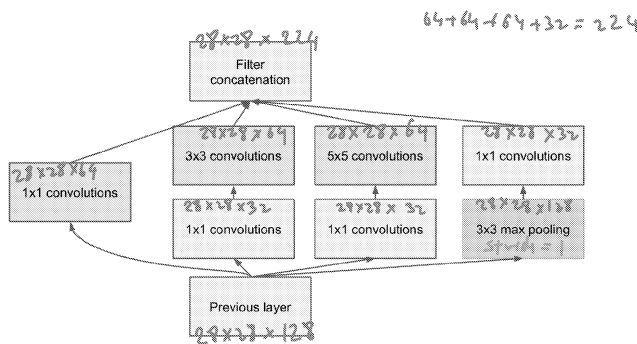


(a) Inception module, naïve version

* multiple receptive fields concatenated \Rightarrow dimension increase

CNN architectures

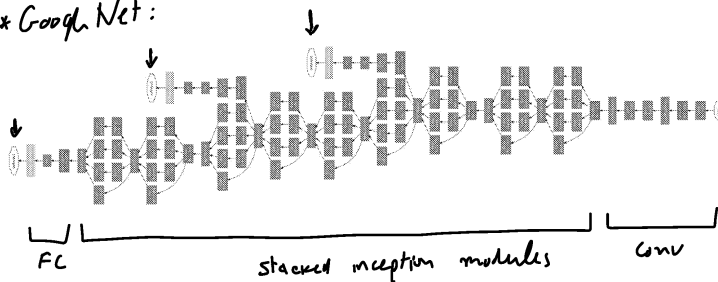
- To reduce the number of parameters use a modified inception block that reduces parameters by using a 1x1 convolution to reduce the number of channels



(b) Inception module with dimensionality reduction

CNN architectures

* Google Net:



- A single fully connected layer results in less parameters
- Convergence may be difficult and so use auxiliary classifier outputs to help with vanishing gradients (not needed when using batch normalization)

CNN architectures

* using residual blocks:




Figure 2. Residual learning: a building block.

- It is easier to learn the $F(x)$ residual compared with $H(x)$ because we need to learn deviation from identity instead of a function
- Skip connections help with vanishing gradients

CNN architectures

VGG-19: 19.6 B FLOPs

Plain 34: 3.6 B FLOPs

ResNet 34: 3.6 B FLOPs




Figure 3. Example network architectures for ImageNet. **Left:** the VGG-19 model [41] (19.6 billion FLOPs) as a reference. **Middle:** a plain network with 34 parameter layers (3.6 billion FLOPs). **Right:** a residual network with 34 parameter layers (3.6 billion FLOPs). The dotted shortcuts increase dimensions. **Table 1** shows more details and other variants.

MobileNets

- **MobileNets:**
 - For mobile and embedded devices
 - Trade-off between latency and accuracy
 - Use depth-wise separable convolutions

Separable convolution. Separate convolution into:


- Depth-wise (channel) convolution followed by
- Point-wise 1×1 convolution

Simplification parameters:

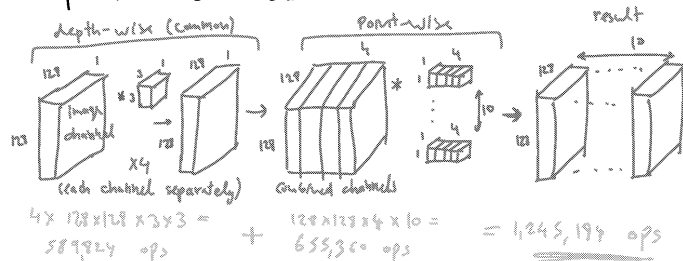
- Width multiplier (fewer channels)
- Resolution multiplier (lower resolution)

MobileNets

* ordinary convolution:



* separable convolution:



MobileNets

* ordinary convolution:

$$D_f \times D_f \times M \xrightarrow{\text{input}} (W \times D_K \times D_K \times M) \xrightarrow{\text{filters}} D_f \times D_f \times N \xrightarrow{\text{output}}$$

$$(D_K \times D_K \times M \times N) \times (D_f \times D_f) \text{ ops}$$

* separable convolution:

$$\begin{aligned} D_f \times D_f \times M &\xrightarrow{\text{input}} (M \times D_K \times D_K \times 1) \xrightarrow{\text{depth-wise}} D_f \times D_f \times M \xrightarrow{\text{depth-wise output}} \\ &\xrightarrow{\text{point-wise}} (N \times 1 \times 1 \times M) \xrightarrow{\text{output}} D_f \times D_f \times N \end{aligned}$$

$$(D_K \times D_K \times 1 \times M) \times (D_f \times D_f) + (1 \times 1 \times M \times N) \times (D_f \times D_f) \text{ ops.}$$

MobileNets

* Reduction in computation:

$$\frac{\text{separable depth-wise conv ops}}{\text{ordinary conv ops}} =$$

$$\begin{aligned} &= \frac{D_K \cdot D_K \cdot M \cdot D_F \cdot D_F + M \cdot N \cdot D_F \cdot D_F}{D_K \cdot D_K \cdot M \cdot N \cdot D_F \cdot D_F} \\ &= \frac{1}{N} + \frac{1}{D_K^2} \end{aligned}$$

MobileNets

* Implementation with Batch Normalization




Figure 3. Left: Standard convolutional layer with batchnorm and ReLU. Right: Depthwise Separable convolutions with Depthwise and Pointwise layers followed by batchnorm and ReLU.

MobileNets

* Width multiplier:

- Reduce # of input and output channels to each layer by a factor of α .

$$\begin{aligned} \text{input channels: } M &\rightarrow \alpha M \\ \text{output channels: } N &\rightarrow \alpha N \end{aligned} \quad \alpha \in [0,1] \quad \text{e.g. } 0.5$$

- New computational cost:

$$\underbrace{D_k \times D_k \times \alpha M \times D_f \times D_f}_{\text{reduction by factor } \alpha} + \underbrace{\alpha M \times \alpha N \times D_f \times D_f}_{\text{reduction by factor } \alpha^2}$$

MobileNets

* Resolution multiplier:

- reduce input resolution by a factor $s \in [0,1]$

$$\begin{aligned} \text{input resolution: } D_f \times D_f &\rightarrow s D_f \times s D_f \\ \text{output resolution: } D_f \times D_f &\rightarrow s D_f \times s D_f \end{aligned}$$

- New computational cost:

$$\underbrace{D_k \times D_k \times \alpha M \times s D_f \times s D_f + \alpha M \times \alpha N \times s D_f \times s D_f}_{\text{reduction by factor } s^2}$$

MobileNets

* Experimental evaluation

Table 3. Resource usage for modifications to standard convolution. Note that each row is a cumulative effect adding on top of the previous row. This example is for an internal MobileNet layer with $D_K = 3$, $M = 512$, $N = 512$, $D_F = 14$.

Layer/Modification	Million	Million
	Mult-Adds	Parameters
Convolution	462	2.36
Depthwise Separable Conv	52.3	0.27
$\alpha = 0.75$	29.6	0.15
$\rho = 0.714$	15.1	0.15

of operations

Table 4. Depthwise Separable vs Full Convolution MobileNet

Model	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
Conv MobileNet	71.7%	4866	29.3
MobileNet	70.6%	569	4.2

accuracy

depth-wise separable conv.

MobileNets

* Experimental evaluation (Contd.)

Table 6. MobileNet Width Multiplier

Width Multiplier	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
1.0 MobileNet-224	70.6%	569	4.2
0.75 MobileNet-224	68.4%	325	2.6
0.5 MobileNet-224	63.7%	149	1.3
0.25 MobileNet-224	50.6%	41	0.5

Table 7. MobileNet Resolution

Resolution	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
1.0 MobileNet-224	70.6%	569	4.2
1.0 MobileNet-192	69.1%	418	4.2
1.0 MobileNet-160	67.2%	290	4.2
1.0 MobileNet-128	64.4%	186	4.2