

Regularization measures

Large number of parameters tend to over-fit
 \Rightarrow need to prevent over-fitting

* Dropout:

- at each training stage drop out units in fully connected layers with probability of $(1-p)$, where p is hyper-parameter
- Removed nodes are reinstated with original weights in the subsequent stage.

Regularization measures

* Data augmentation:

- increase variability in training data
- Perturb existing data (e.g. crop images in different ways)

* Early stopping:

- stop training before learning completes

Batch normalization

- Makes sure that activations are not saturated
- Normalization values are computed for each batch in training
- Normalization is differentiable (suitable for backpropagation)
- In addition to normalization allow for some shift and scaling to support some saturation and end training
- Batch normalization is not needed after each layer

batch outputs $\{z^{(i)}\}_{i=1}^q \rightarrow \{\hat{z}^{(i)}\}_{i=1}^q$
 for current layer

$$\hat{z}_j^{(i)} = \frac{z_j^{(i)} - \mu_j}{\sigma_j}$$


\uparrow
 output of j -th unit
 for i -th batch example

$$\mu_j = \frac{1}{q} \sum_{i=1}^q z_j^{(i)}$$


$$\sigma_j = \left(\frac{1}{q} \sum_{i=1}^q (z_j^{(i)} - \mu_j)^2 \right)^{1/2}$$

Batch normalization

* Usually normalize before activation:



* Layers:



FC = Fully connected or convolutional layer without activation

BN = Batch normalization

σ = NL activation

* sometimes applied after activation.

Ensemble classifiers

* Ensemble classifier:

- Train multiple independent models
- use majority vote or average during testing

* Reduces overfitting

* To obtain multiple models:

- change data
- change parameters
- Record multiple snapshots of the model during training.

Summary

* start without regularization

* add Batch normalization

* add dropout if needed

* add decay if needed

* augment data if needed

* create ensemble if needed

if observing overfitting



Hyperparameter optimization

* principle:

- select a set of parameters
- train on training data and test on validation data using a few epochs
- modify parameters in search of better validation error.
- If loss is increasing (e.g. $\times 2$) stop

Hyperparameter optimization

* Methods:

- Coordinate search
- Grid search
- Random search
- Bayesian model-based optimization
- Evolutionary

Hyperparameter optimization


* Search procedures

- hold all parameters fixed and change one
- change parameter on a logarithmic scale to cover both small and large values (e.g. $\times 2$ or $\times 10$ values)
- once identifying best parameter search around it (coarse-to-fine search) using more epochs
- If best parameters are found at end of range extend range.
- Repeat several passes on all parameters


Hyperparameter optimization

* Grid and random search

- Search combinations of parameters:



Hyperparameter optimization



quasi-random sampling
(low discrepancy sequence)

random sampling

⇒ Better coverage

Hyperparameter optimization

* Optimize parameters in this order:

- learning rate (value, decay, update procedure)
- regularization (batch norm, dropout, L2)
- network architecture and network size
- other:
 - optimizer
 - optimizer parameters
 - activation
 - initialization

GPU frameworks

* Framing is time consuming \Rightarrow use GPUs

* Frameworks:

Tensorflow

Keray

Pytorch

[illegible]